

Lab 5: Investigating the Performance Advantages of NumPy Arrays

At a high level, the purpose of this lab is to test how much more efficient bulk, vectorized operations can be (using NumPy arrays) than standard, dictionary-oriented Python computations.

In this lab, your task is to compare three different solutions to the *co-occurrence problem*. This is a very fundamental computation in text analytics. In a nutshell, we have n documents. We have a dictionary with some number of words; we know how many times each document has each word in the dictionary. For each of the possible (word, word) pairs, we want to compute the number of documents (between zero and n) that have that particular (word, word) pair.

For example, let's say our documents are:

```
doc 1: {word1, word2, word4, word5}
doc 2: {word1, word2, word5}
doc 3: {word2, word3, word5}
```

Then the result of the co-occurrence computation is:

```
{word1, word1}: 2 co-occurs (means word1 occurs twice in the corpus)
{word1, word2}: 2 co-occurs
{word1, word4}: 1 co-occurs
{word1, word5}: 2 co-occurs
{word2, word2}: 3 co-occurs
{word2, word3}: 1 co-occurs
{word2, word4}: 1 co-occurs
{word2, word5}: 3 co-occurs
{word3, word3}: 1 co-occurs
{word3, word5}: 1 co-occurs
{word4, word4}: 1 co-occurs
{word4, word5}: 1 co-occurs
{word5, word5}: 3 co-occurs
```

Given that, here are the lab's three subtasks:

Subtask One

First, run the pure, dictionary-based LDA implementation provided (this is the first one) to build a document corpus (just like in the LDA lab). This will build the `wordsInCorpus` object, which is a Python dictionary. The dictionary's key is a document identifier, and the value is another dictionary. For the dictionary associated with a particular document identifier, the key is a word identifier, and the value is the number of occurrences of the word in the document.

Given this setup, write and execute the following code:

```
start = time.time()

# coOccurrences will be a map where the key is a
# (wordOne, wordTwo) pair, and the value is the number of times
# those two words co-occurred in a document, so this will be a
# value between zero and 50
coOccurrences = {}

# now, have a nested loop that fills up coOccurrences
# YOUR CODE HERE

end = time.time()
end - start
```

You should replace the `# YOUR CODE HERE` with a nested loop that loops through all documents. For each document, it loops through all pairs of words in the document, and increments a count for the pair. The count for all pairs should be put into a dictionary where the key is a tuple of the form `(word1, word2)` and the value is a count.

In case you don't remember, the way to loop through all of the keys in a Python dictionary is something like the following:

```
for wordOne in wordsInCorpus[doc]:
    # code here
```

Subtask Two

There is another way to compute the same result, this time using NumPy arrays. Note that outer product of a vector x with itself creates a matrix where the entry at row i , column j is $x[i] * x[j]$. Let's say that we have a vector, where the i th entry in the vector is 1 if the i th word in the dictionary appears in the document, and a 0 otherwise. Then, if we take an outer product of this vector with itself, we'll have a matrix where the entry at row i , column j is 1 if the document has word i and word j .

Given this, write a NumPy code that loops through all of the documents. Treat each document as a vector whose length is the length of the corpus. Then you can use the outer product of the vector giving the word counts for that document with itself to create a matrix of co-occurrences for that document. Summing the 50 matrices gives the answer. Hint: in order to take a NumPy array `foo` and "clip" all of its entries so that none of them is greater than one, use the code `np.clip (foo, 0, 1)`. And `np.outer (foo, bar)` computes the outer product of two vectors/matrices.

To complete this task, first run the NumPy array-based LDA implementation (again, see

Canvas). This will build the `wordsInCorpus` array. It stores the same data, but in a NumPy array. `wordsInCorpus[34, 355]` is the number of times that the 355th word in the dictionary occurred in the 34th document.

And Here is the skeleton for the code you will write:

```
start = time.time()

# coOccurrences[i, j] will give the count of the number of times that
# word i and word j appear in the same document in the corpus
coOccurrences = np.zeros ((2000, 2000))

# now, have a nested loop that fills up coOccurrences
# YOUR CODE HERE

end = time.time()
end - start
```

Subtask Three

We could do essentially the same thing with just matrix multiply. Now, write a code that uses a single matrix multiply to create `coOccurrences`. Note that you can use the `transpose` operation to transpose a matrix, and `np.dot ()` will multiply two matrices.

```
start = time.time()

# now, create coOccurrences via a matrix multiply
# YOUR CODE HERE

end = time.time()
end - start
```

What to Turn In

In Canvas, simply submit the output you get from running each of your three codes, as well as the codes themselves (just copy and paste into Canvas, or put together a single .txt file or .pdf file, and submit that).